

Exploratory Data Analysis

June 20, 2021

1 ANZ Virtual Internship Report

Task 1 Exploratory data analysis

```
In [1]: import os
        cwd = os.getcwd()
        cwd
```

```
Out[1]: 'C:\\\\Users\\hp'
```

```
In [2]: os.chdir("C:/Users/hp/Desktop/Data Science-Data@ANZ/Module 1")
```

```
In [3]: os.listdir('.')
```

```
Out[3]: ['ANZ synthesised transaction dataset.xlsx',
        'anz_cleaned_data.xlsx',
        'Data@ANZ.ipynb',
        'Model Answer for Task 1.pdf',
        '~$ANZ synthesised transaction dataset.xlsx']
```

```
In [4]: import xlrd
        import numpy as np
        import pandas as pd
        import datetime as dt
        import seaborn as sns
        import matplotlib.pyplot as plt
```

```
In [5]: file = "ANZ synthesised transaction dataset.xlsx"
        xl = pd.ExcelFile(file)
        print(xl.sheet_names)
```

```
['DSynth_Output_100c_3m_v3']
```

```
In [6]: df = pd.read_excel(file, sheet_name = "DSynth_Output_100c_3m_v3")
        df.head()
```

```

Out[6]:      status  card_present_flag  bpay_biller_code      account  currency  \
0  authorized                1.0                NaN  ACC-1598451071    AUD
1  authorized                0.0                NaN  ACC-1598451071    AUD
2  authorized                1.0                NaN  ACC-1222300524    AUD
3  authorized                1.0                NaN  ACC-1037050564    AUD
4  authorized                1.0                NaN  ACC-1598451071    AUD

      long_lat  txn_description      merchant_id  \
0  153.41 -27.95          POS  81c48296-73be-44a7-befa-d053f48ce7cd
1  153.41 -27.95      SALES-POS  830a451c-316e-4a6a-bf25-e37caedca49e
2  151.23 -33.94          POS  835c231d-8cdf-4e96-859d-e9d571760cf0
3  153.10 -27.66      SALES-POS  48514682-c78a-4a88-b0da-2d6302e64673
4  153.41 -27.95      SALES-POS  b4e02c10-0852-4273-b8fd-7b3395e32eb0

      merchant_code  first_name  ...      age  merchant_suburb  merchant_state  \
0              NaN      Diana  ...      26      Ashmore          QLD
1              NaN      Diana  ...      26      Sydney          NSW
2              NaN    Michael  ...      38      Sydney          NSW
3              NaN      Rhonda  ...      40      Buderim          QLD
4              NaN      Diana  ...      26  Mermaid Beach          QLD

      extraction  amount      transaction_id  \
0  2018-08-01T01:01:15.000+0000  16.25  a623070bfead4541a6b0fff8a09e706c
1  2018-08-01T01:13:45.000+0000  14.19  13270a2a902145da9db4c951e04b51b9
2  2018-08-01T01:26:15.000+0000   6.42  feb79e7ecd7048a5a36ec889d1a94270
3  2018-08-01T01:38:45.000+0000  40.90  2698170da3704fd981b15e64a006079e
4  2018-08-01T01:51:15.000+0000   3.25  329adf79878c4cf0aeb4188b4691c266

      country  customer_id  merchant_long_lat  movement
0  Australia  CUS-2487424745      153.38 -27.99  debit
1  Australia  CUS-2487424745      151.21 -33.87  debit
2  Australia  CUS-2142601169      151.21 -33.87  debit
3  Australia  CUS-1614226872      153.05 -26.68  debit
4  Australia  CUS-2487424745      153.44 -28.06  debit

[5 rows x 23 columns]

```

2 Data preparation

The dataset contains 12043 transactions for 100 customers who have one bank account each. Transactional period is from 01/08/2018 - 31/10/2018 (92 days duration). The data entries are unique and have consistent formats for analysis. For each record/row, information is complete for majority of columns. Some columns contain missing data (blank or NA cells), which is likely due to the nature of transaction. (i.e. merchants are not involved for InterBank transfers or Salary payments) It is also noticed that there is only 91 unique dates in the dataset, suggesting the transaction records for one day are missing (turned out to be 2018-08-16). The range of each feature should also be examined which shows that there is one customer that resides outside Australia

```
In [7]: df.shape
```

```
Out[7]: (12043, 23)
```

```
In [8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12043 entries, 0 to 12042
Data columns (total 23 columns):
status                12043 non-null object
card_present_flag     7717 non-null float64
bpay_biller_code      885 non-null object
account              12043 non-null object
currency              12043 non-null object
long_lat              12043 non-null object
txn_description        12043 non-null object
merchant_id           7717 non-null object
merchant_code         883 non-null float64
first_name            12043 non-null object
balance               12043 non-null float64
date                  12043 non-null datetime64[ns]
gender                12043 non-null object
age                   12043 non-null int64
merchant_suburb       7717 non-null object
merchant_state        7717 non-null object
extraction            12043 non-null object
amount                12043 non-null float64
transaction_id        12043 non-null object
country               12043 non-null object
customer_id           12043 non-null object
merchant_long_lat     7717 non-null object
movement              12043 non-null object
dtypes: datetime64[ns](1), float64(4), int64(1), object(17)
memory usage: 2.1+ MB
```

3 Data Cleaning Stage

```
In [9]: # checking the unique values in currency
df.currency.value_counts()
```

```
Out[9]: AUD      12043
Name: currency, dtype: int64
```

```
In [10]: # dropping currency column as of no use in analysis has no unique value.
df.drop(['currency'], axis = 1 , inplace = True)
```

```
In [11]: # checking for nan in dataset
df.isnull().sum()
```

```
Out[11]: status          0
        card_present_flag 4326
        bpay_biller_code 11158
        account          0
        long_lat         0
        txn_description   0
        merchant_id      4326
        merchant_code     11160
        first_name        0
        balance          0
        date             0
        gender            0
        age              0
        merchant_suburb   4326
        merchant_state    4326
        extraction       0
        amount           0
        transaction_id    0
        country           0
        customer_id       0
        merchant_long_lat 4326
        movement         0
        dtype: int64
```

```
In [12]: df.drop(['bpay_biller_code', 'merchant_code'], axis = 1 , inplace = True)
```

```
In [13]: # count# these are the unique values in the status column.
        df['status'].value_counts()
        df.duplicated().sum()
```

```
Out[13]: 0
```

```
In [14]: # it looks like the NaNs in the dataset are all on common rows.
        df[df.card_present_flag.isnull()][['merchant_id', 'merchant_state', 'merchant_suburb']
```

```
Out[14]: merchant_id      4326
        merchant_state    4326
        merchant_suburb   4326
        merchant_long_lat 4326
        dtype: int64
```

```
In [15]: df['status'].value_counts()
```

```
Out[15]: authorized      7717
        posted          4326
        Name: status, dtype: int64
```

```
In [16]: # all the transactions happened in Australia.
        df['country'].value_counts()
        # then we can drop the country column
        df.drop(['country'], axis = 1 , inplace = True)
```

```

In [17]: # the unique movement types
df.movement.value_counts()

Out[17]: debit      11160
         credit      883
         Name: movement, dtype: int64

In [18]: # dtype of date column and extraction column to datetime
df['date'] = pd.to_datetime(df['date'])
df['extraction'] = pd.to_datetime(df['extraction'])

In [19]: # adding three columns: year, month, and day of the transaction to make the segmentat
df['day']= df['date'].dt.day_name()
df['month'] = df['date'].dt.month_name()
df['year'] = df['date'].dt.year
df['hour']= df.extraction.dt.hour

In [20]: # unique values of our new columns
df.year.value_counts()

Out[20]: 2018      12043
         Name: year, dtype: int64

In [21]: # because all the transactions are from 2018, I will drop the year columns
df.drop(['year'], axis = 1 , inplace = True)

In [22]: # it looks like transactions happened only in August, September, and October
df.month.value_counts()

Out[22]: October      4087
         September    4013
         August       3943
         Name: month, dtype: int64

In [23]: # unique values of txn_description (mode of transaction)
df.txn_description.value_counts()

Out[23]: SALES-POS      3934
         POS           3783
         PAYMENT       2600
         PAY/SALARY     883
         INTER BANK    742
         PHONE BANK    101
         Name: txn_description, dtype: int64

In [24]: # Check the distribution of the dataset by gender
df.gender.value_counts()

Out[24]: M      6285
         F      5758
         Name: gender, dtype: int64

```

```
In [25]: # this is a quick summary of our categorical variables showing the number of unique v
# column with the most frequent value in it and its frequency.
df.describe(exclude = [np.number]).loc[['unique','top','freq']]
```

```
Out [25]:
```

	status	account	long_lat	txn_description	\
unique	2	100	100	6	
top	authorized	ACC-1598451071	153.41 -27.95	SALES-POS	
freq	7717	578	578	3934	

	merchant_id	first_name	date	\
unique	5725	80	91	
top	106e1272-44ab-4dcb-a438-dd98e0071e51	Michael	2018-09-28 00:00:00	
freq	14	746	174	

	gender	merchant_suburb	merchant_state	extraction	\
unique	2	1609	8	9442	
top	M	Melbourne	NSW	2018-09-03 11:00:00	
freq	6285	255	2169	10	

	transaction_id	customer_id	merchant_long_lat	\
unique	12043	100	2703	
top	aa78c28d430240d49bfed5aa4a8bb42f	CUS-2487424745	151.21 -33.87	
freq	1	578	145	

	movement	day	month
unique	2	7	3
top	debit	Friday	October
freq	11160	2073	4087

transaction_id has 12043 unique values. it won't be significant in my analysis so I will drop it. there is 100 customers in the dataset. However, there are only 80 unique first names. I decided to drop the first name columns because the customer_id is more accurate. because each customer has a unique account. I will only keep the customer_id column and drop account column. for the location based analysis, I will rely merchant_state. I will drop long_lat, merchant_suburb, and merchant_long_lat. there are 5725 merchants. the highest number of transactions by merchant is 14 only, meaning the merchant_id won't significantly have an impact on the amount of the transaction. I will drop for now.

```
In [26]: df.drop(['transaction_id', 'account', 'first_name', 'long_lat', 'merchant_id', 'merchant
```

```
In [27]: #final dataset
```

```
df.head()
```

```
Out [27]:
```

	status	card_present_flag	txn_description	balance	date	gender	\
0	authorized	1.0	POS	35.39	2018-08-01	F	
1	authorized	0.0	SALES-POS	21.20	2018-08-01	F	
2	authorized	1.0	POS	5.71	2018-08-01	M	
3	authorized	1.0	SALES-POS	2117.22	2018-08-01	F	

```

4   authorized          1.0      SALES-POS    17.95 2018-08-01      F

      age merchant_state      extraction  amount  customer_id movement \
0    26                QLD 2018-08-01 01:01:15   16.25 CUS-2487424745  debit
1    26                NSW 2018-08-01 01:13:45   14.19 CUS-2487424745  debit
2    38                NSW 2018-08-01 01:26:15    6.42 CUS-2142601169  debit
3    40                QLD 2018-08-01 01:38:45   40.90 CUS-1614226872  debit
4    26                QLD 2018-08-01 01:51:15    3.25 CUS-2487424745  debit

      day  month  hour
0  Wednesday  August    1
1  Wednesday  August    1
2  Wednesday  August    1
3  Wednesday  August    1
4  Wednesday  August    1

```


Univariate analysis First, a univariate analysis for my categorical variables.

In [28]: *# there are more authorized transactions than posted*

```

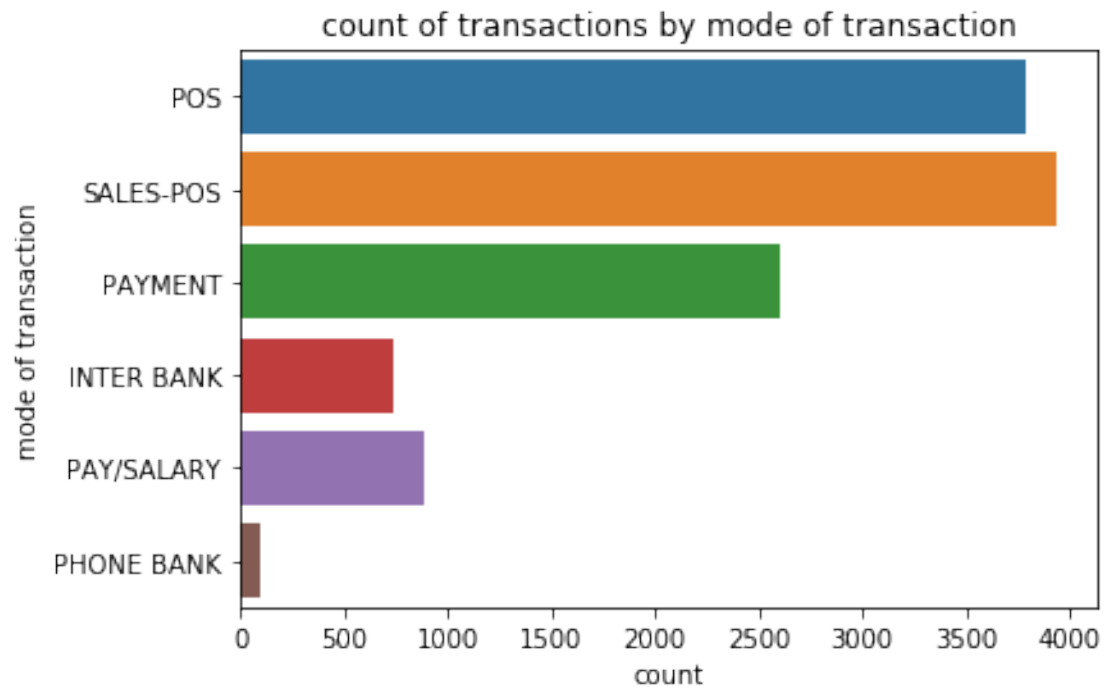
sns.countplot(x = 'status', data = df)
plt.title('count of transactions by transaction status');

```




In [29]: *# POS and SALES-POS are the most used modes of transaction*

```
sns.countplot(y = 'txn_description', data = df)
plt.title('count of transactions by mode of transaction')
plt.ylabel('mode of transaction');
```




In [30]: *# there are more males than females*

```
sns.countplot(x = 'gender', data = df)
plt.title('count of transactions by gender');
```


```
In [31]: # NSW and VIC are the states with the most transactions
```

```
sns.countplot(x = 'merchant_state', data = df)  
plt.title('count of transactions by merchants state');
```



In [32]: *# there are much more debit transactions than credit ones.*

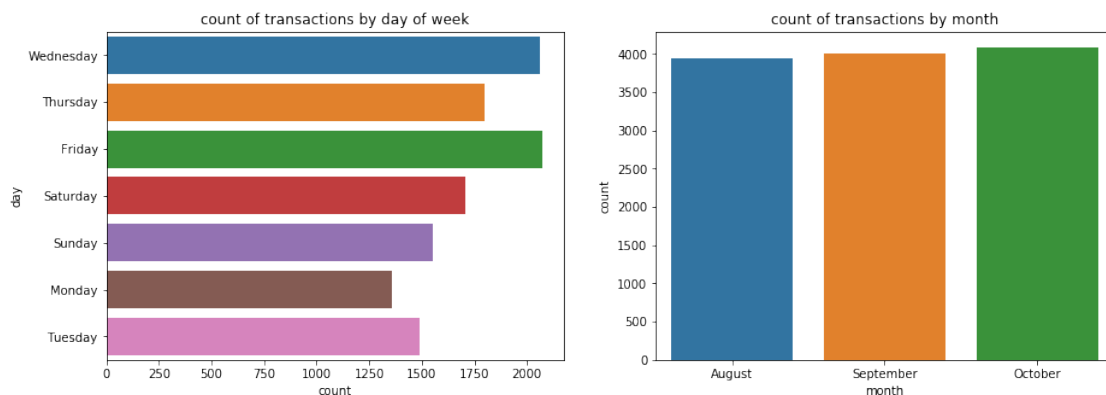
```
sns.countplot(x = 'movement', data = df)
plt.title('count of transactions by movement type');
```



```
In [33]: plt.figure(figsize=[15, 5])

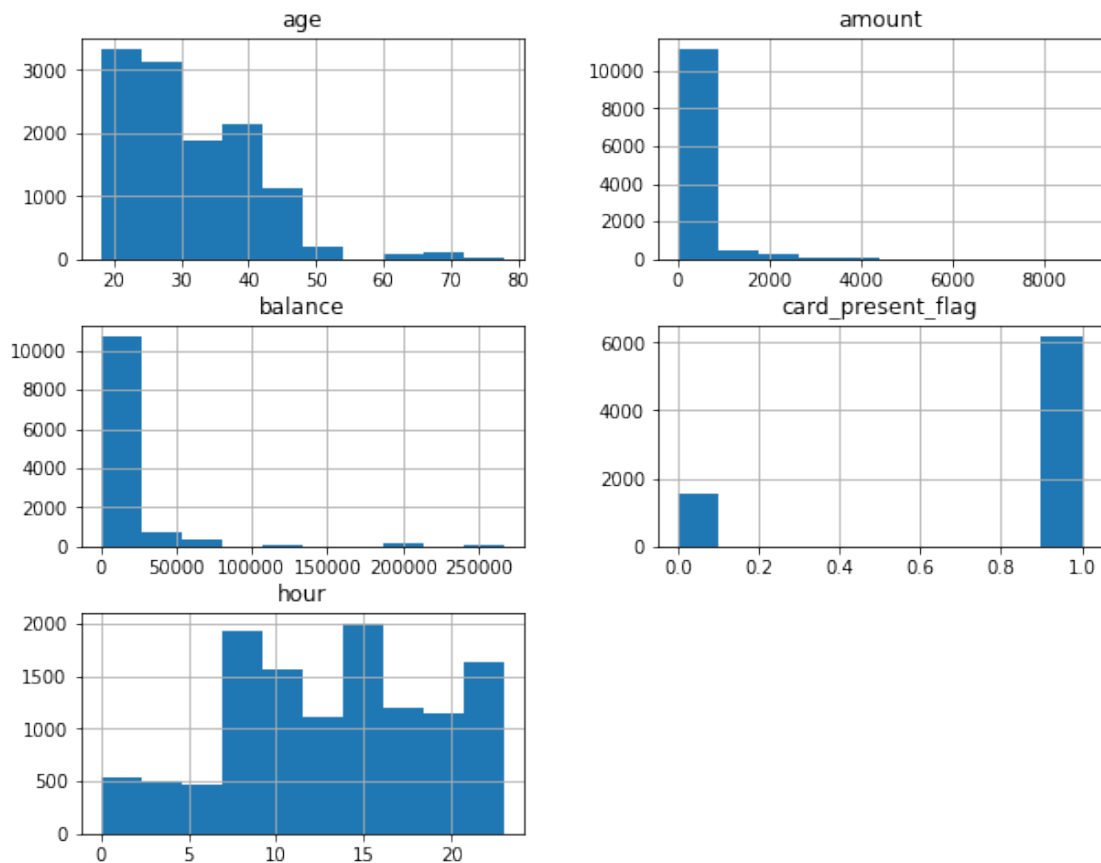
# the left plot
plt.subplot(1,2,1)
sns.countplot(y = 'day', data = df)
plt.title('count of transactions by day of week')

# the right plot
plt.subplot(1,2,2)
sns.countplot(x = 'month', data = df)
plt.title('count of transactions by month');
```



Wednesday and Friday are the days with the most number of transactions. Monday scored the lowest number of transactions the difference in the number of transactions between the months is not significant

In [34]: *# these are histograms of my numeric variables.*
`df.hist(figsize=(10,8));`



In [35]: `df.describe(percentiles = [0.75, 0.90, 0.95, 0.99])`

```
Out [35]:
```

	card_present_flag	balance	age	amount \
count	7717.000000	12043.000000	12043.000000	12043.000000
mean	0.802644	14704.195553	30.582330	187.933588
std	0.398029	31503.722652	10.046343	592.599934
min	0.000000	0.240000	18.000000	0.100000
50%	1.000000	6432.010000	28.000000	29.000000
75%	1.000000	12465.945000	38.000000	53.655000
90%	1.000000	29442.384000	43.000000	186.000000
95%	1.000000	53362.930000	46.000000	1158.150000


99%	1.000000	201963.445200	69.000000	3195.010000
max	1.000000	267128.520000	78.000000	8835.980000

	hour
count	12043.000000
mean	13.268621
std	5.777284
min	0.000000
50%	13.000000
75%	18.000000
90%	21.000000
95%	22.000000
99%	23.000000
max	23.000000

There seem to be some outliers in amount, age, and balance columns.

4 Transaction amount analysis


```
In [36]: plt.figure(figsize=[15,4])
        binsize = 10
        bins = np.arange(-100, df['amount'].max()+binsize, binsize)
        plt.hist(data = df, x = 'amount', bins = bins)
        plt.xlabel('transaction amount');
```



Getting rid of the outliers by limiting the x axis

```
In [37]: plt.figure(figsize=[8, 5])

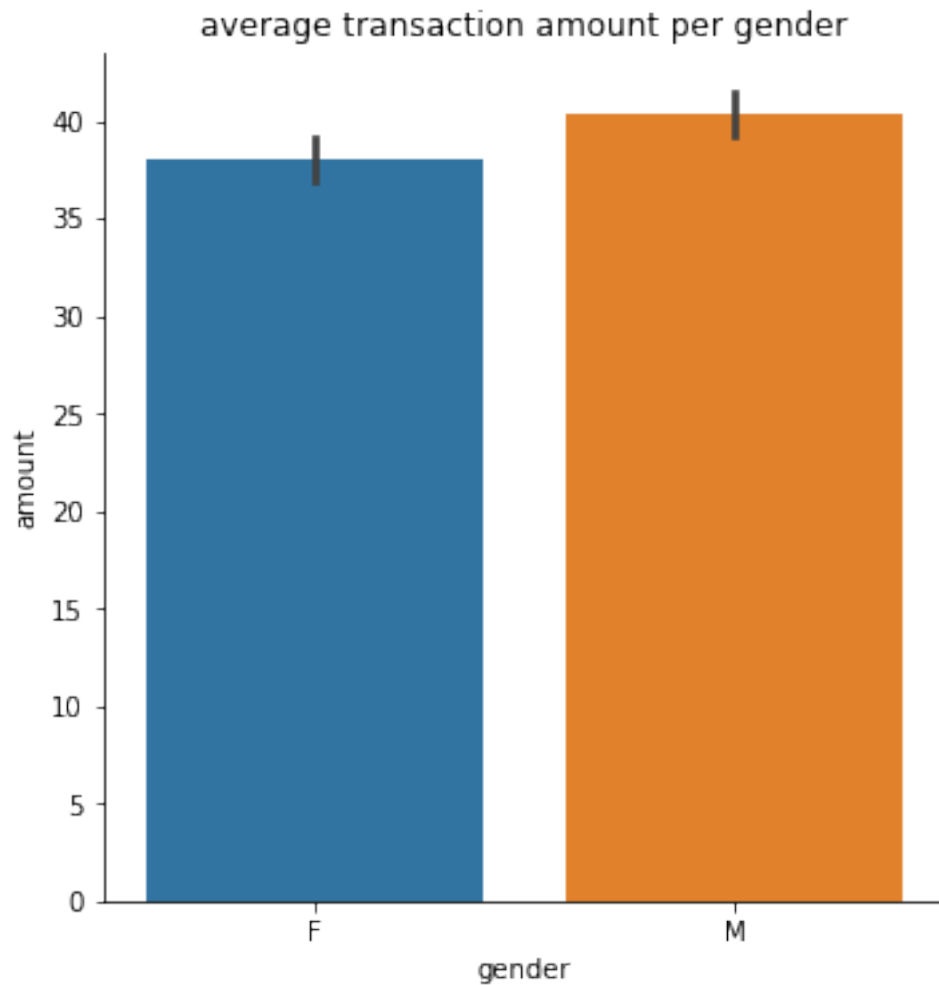
        bin_edges = np.arange(-10, 350+1,1)
        plt.hist(data = df, x = 'amount', bins = bin_edges)
        plt.xlim(-10, 350)
        plt.xlabel('transaction amount');
```



the transaction amount looks normally distributed but long-tailed due to outliers. But with the use of axis limits, we are able to see the normal distribution clearly

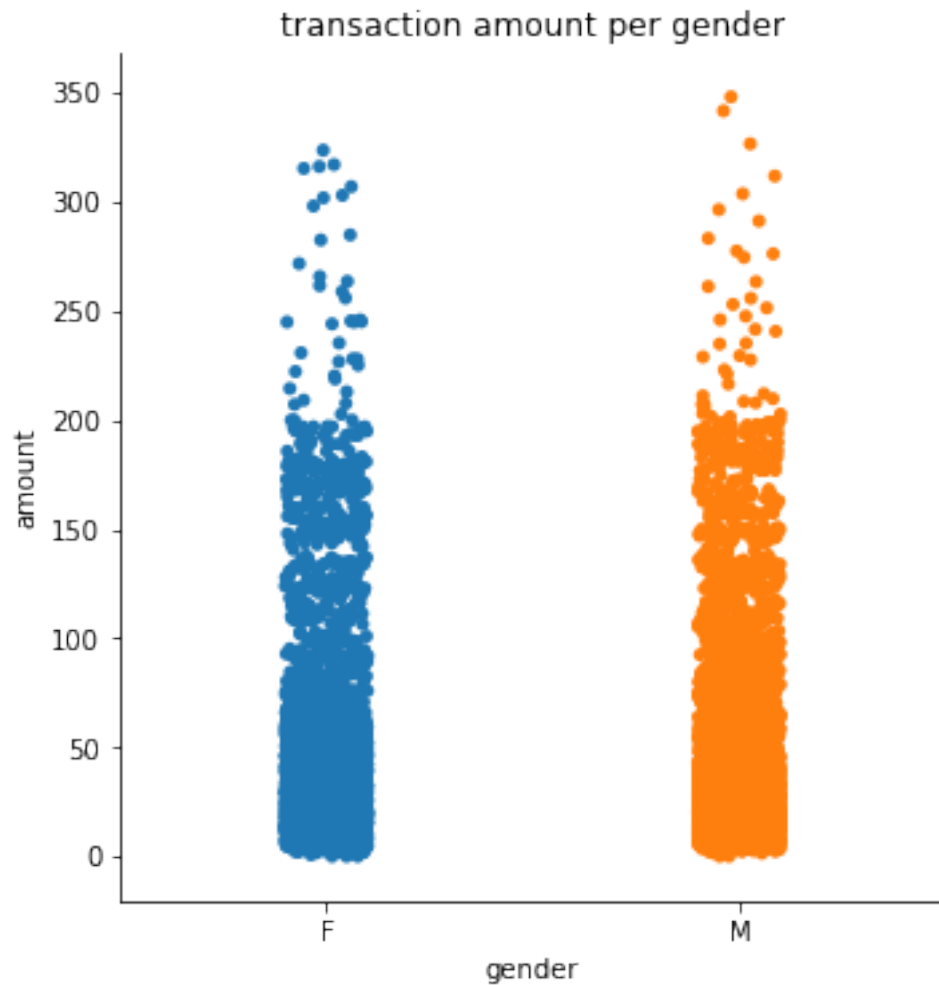
Amount by gender

```
In [38]: sns.catplot(x="gender", y="amount", kind="bar", data=df.query('amount <= 350'))  
         plt.title('average transaction amount per gender');
```



this is a barplot showing the average amount of transactions for each gender. with a confidence interval displayed as error bars.

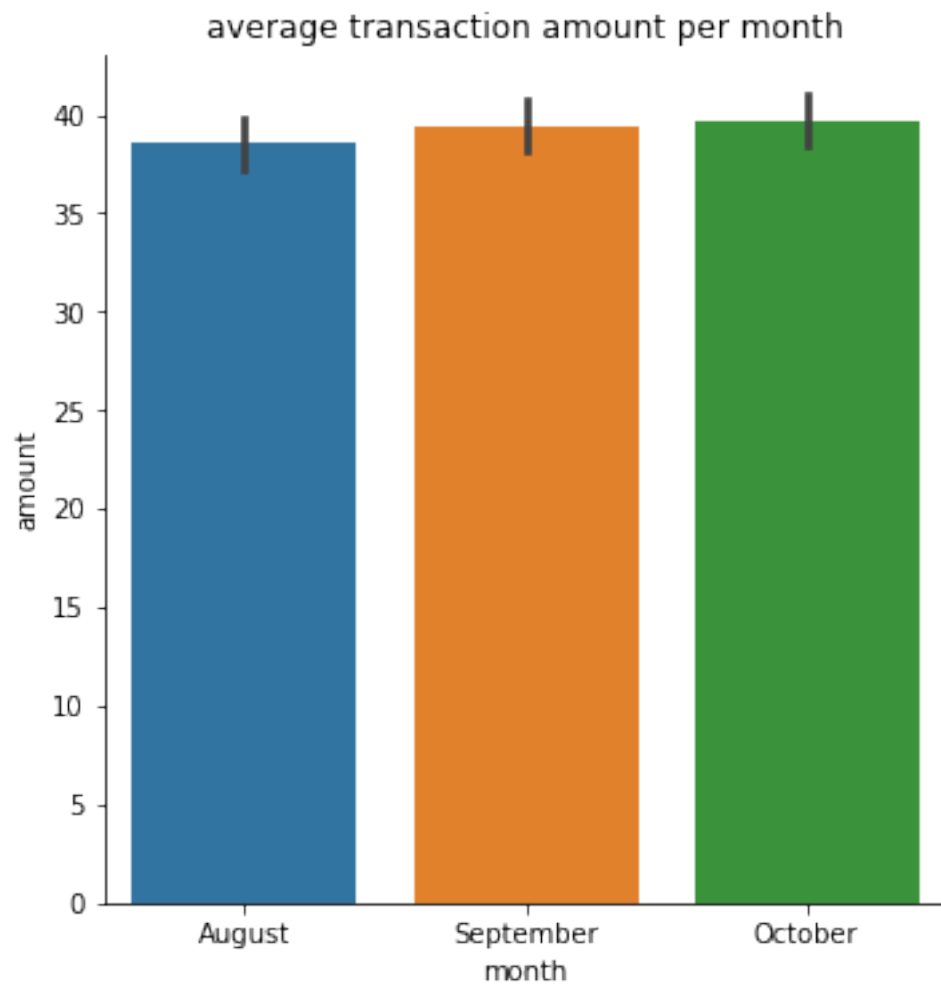
```
In [39]: sns.catplot(x="gender", y = 'amount', data=df.query('amount <= 350'))  
         plt.title('transaction amount per gender');
```




this is a categorical scatterplot showing the transaction amount for each gender.
the highest amounts are made by men the amount of transactions made by men are higher on average.

5 Amount by month


```
In [40]: sns.catplot(x="month", y="amount", kind="bar", data=df.query('amount <= 350'))  
plt.title('average transaction amount per month')  
sns.catplot(x="month", y="amount", kind="point", data=df.query('amount <= 350'));
```




October has higher transaction amounts than August and September.

```
In [41]: sns.catplot(x="month", y = 'amount', data=df.query('amount <= 350'))  
         plt.title('transaction amount per gender');
```



outliers in the transaction amount happen in all three months.

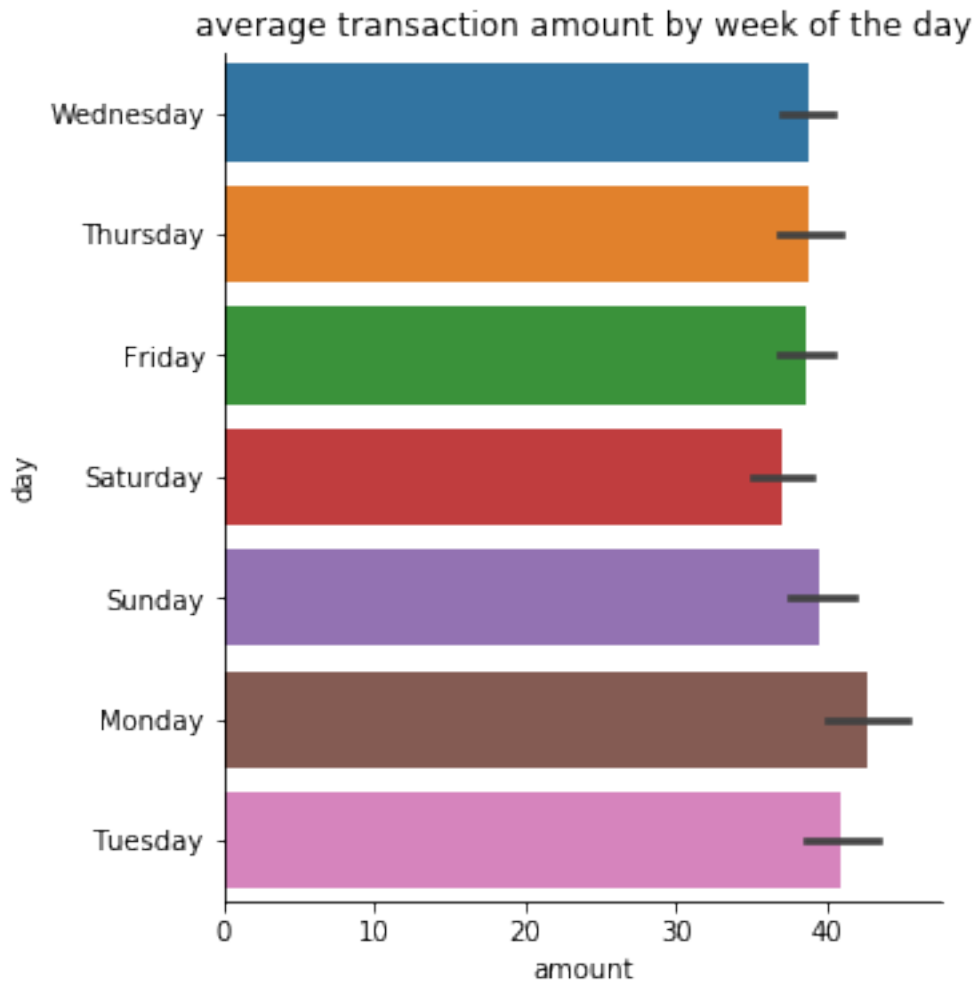
```
In [42]: sns.catplot(x="month", y="amount", hue="gender", kind="point", data=df.query('amount < 300'))  
plt.title('average transaction amount by month per gender');
```




males have higher transaction amounts. October is the month with the highest transaction amounts for both genders.

6 Amount by day


```
In [43]: sns.catplot(y="day", x="amount", kind="bar", data=df.query('amount <= 350'))
plt.title('average transaction amount by week of the day')
sns.catplot(y="day", x="amount", kind="point", data=df.query('amount <= 350'));
```






the amount of transactions is the lowest on weekend Saturday. the transactions with the highest amounts happen mostly on Monday.

```
In [44]: sns.catplot(y="day", x="amount", hue="gender", kind="point", data=df.query('amount <= 50000'))  
plt.title('average transaction amount by weekday per gender');
```



only on Sunday do women have almost the same amount of transaction as men.

```
In [45]: sns.catplot(y="day", x="amount", hue='month', kind="bar", data=df.query('amount <= 350'))
plt.title('transaction amount by week of the day per month');
```



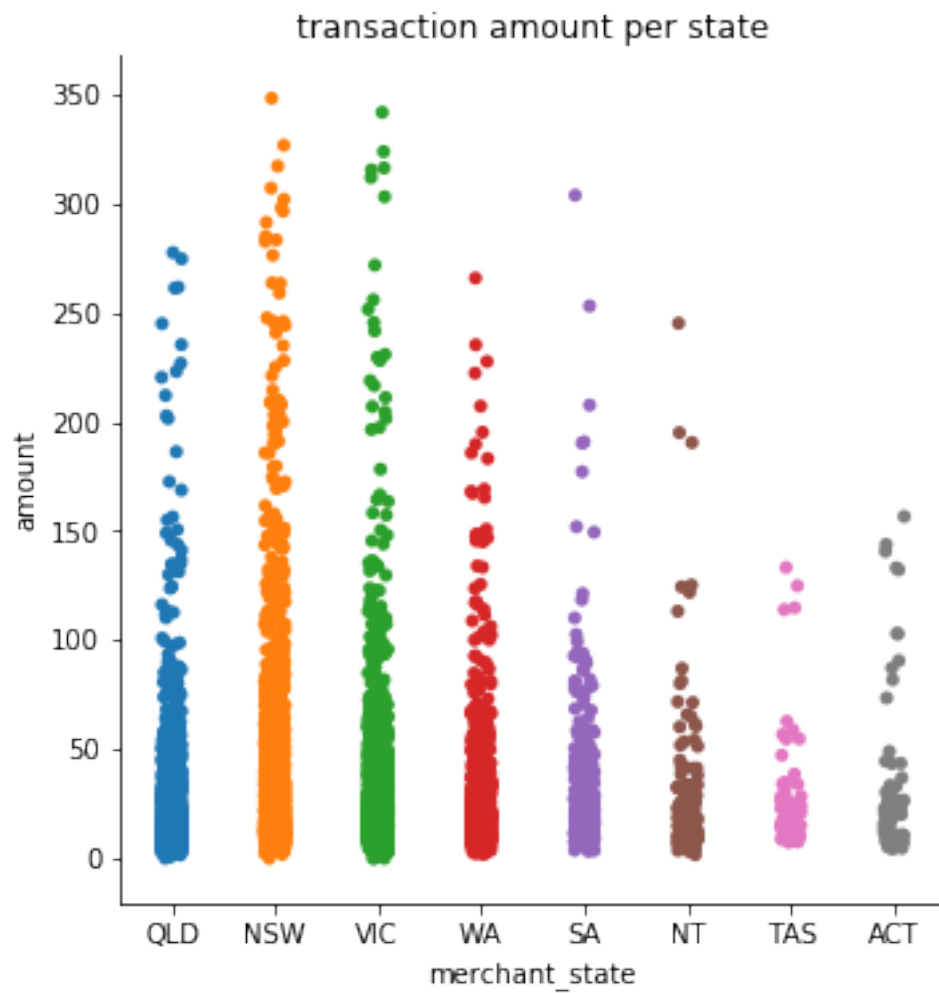
Amount by location

```
In [46]: # min and max amount of transactions in each state
df.groupby('merchant_state')['amount'].agg(['min' , 'max'])
```

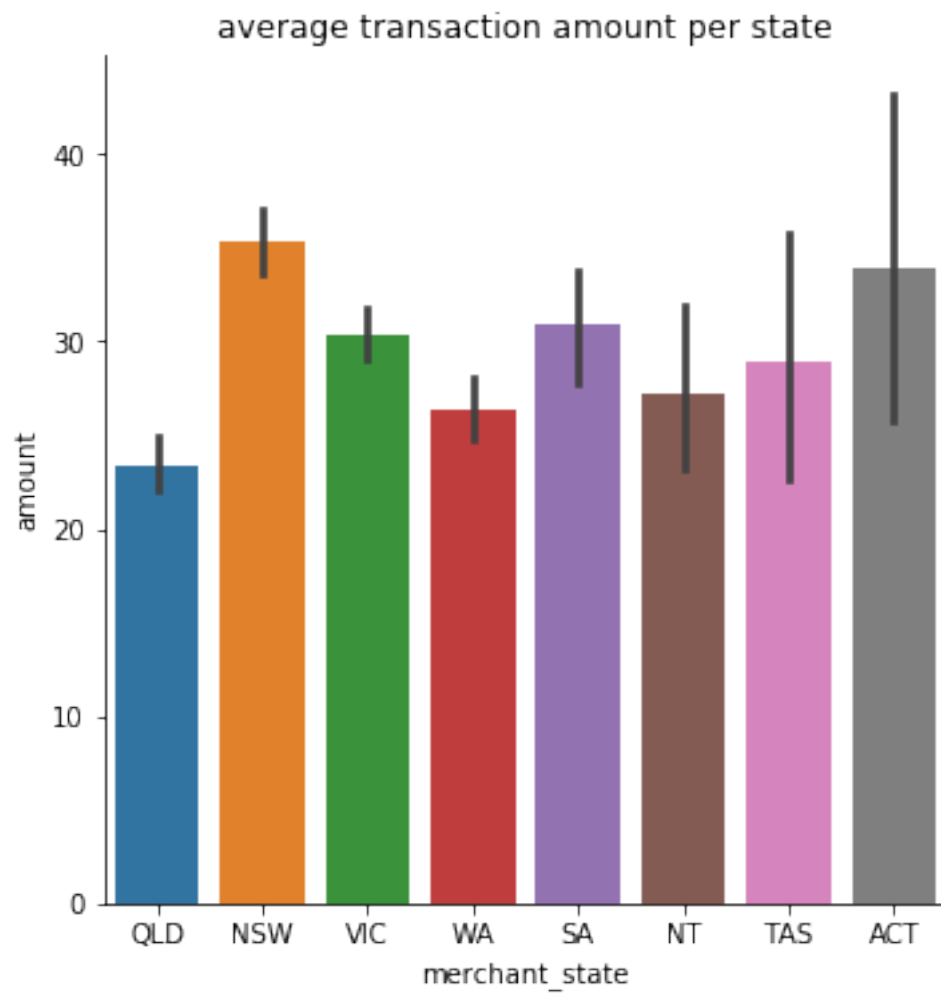
```
Out[46]:
```

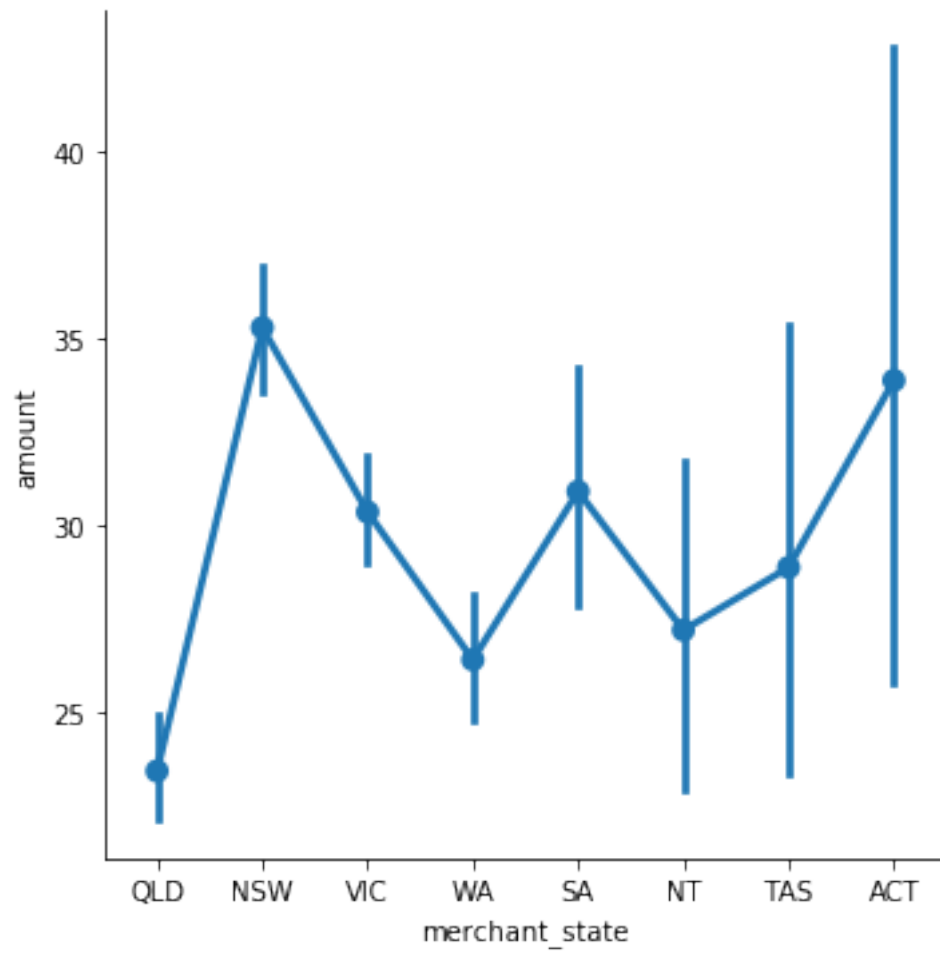
merchant_state	min	max
ACT	4.50	1348.47
NSW	0.10	4233.00
NT	1.71	1285.09
QLD	0.10	7081.09
SA	3.28	909.61
TAS	7.35	133.31
VIC	0.10	3680.71
WA	2.05	1692.56

```
In [47]: sns.catplot(x="merchant_state", y = 'amount', data=df.query('amount <= 350'))
plt.title('transaction amount per state');
```





```
In [48]: sns.catplot(x="merchant_state", y="amount", kind="bar", data=df.query('amount <= 350'))  
plt.title('average transaction amount per state')  
sns.catplot(x="merchant_state", y="amount", kind="point", data=df.query('amount <= 350'))
```





```
In [49]: sns.catplot(y="merchant_state", x="amount", hue="gender", kind="point", data=df.query(
plt.title('average transaction amount by state per gender');
```




women's transaction amounts are higher than men in the state of TAS and ACT


7 Amount by Age

```
In [50]: sns.distplot(df['age']);
```

C:\Users\hp\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is deprecated and will be removed in a future version. Use `displot` instead.
warnings.warn(msg, FutureWarning)



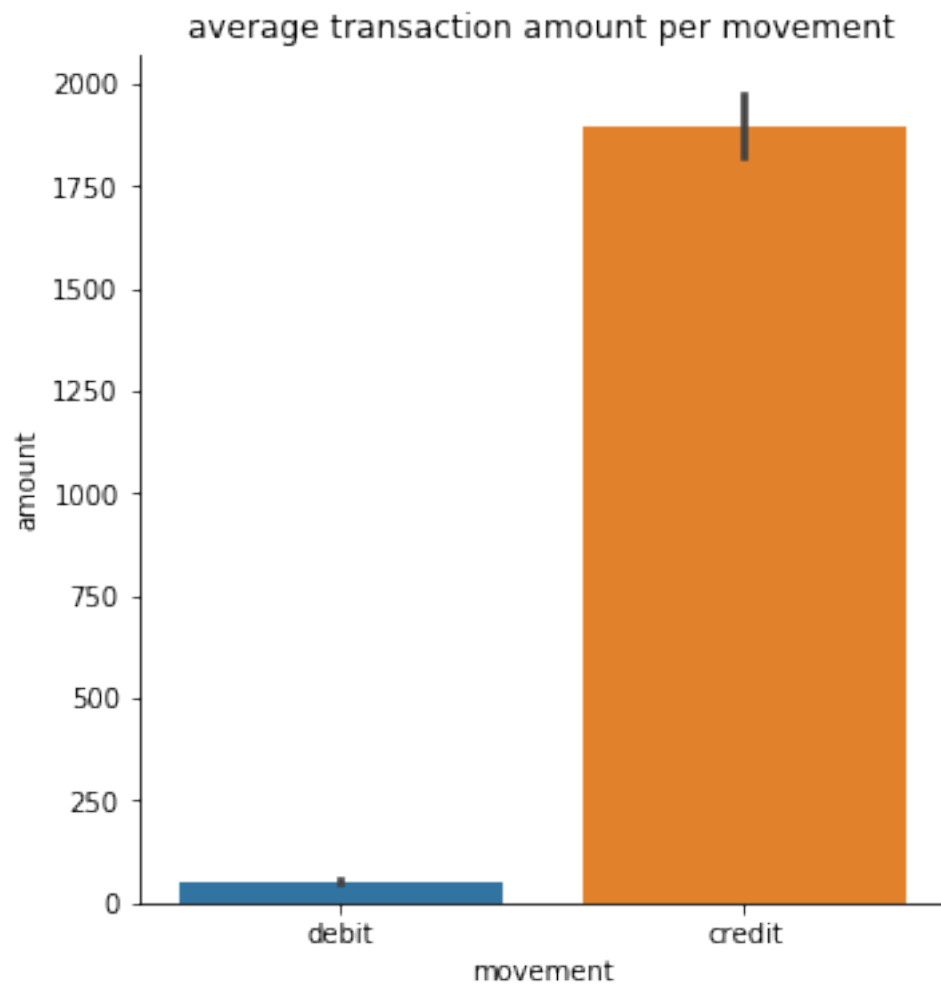
```
In [51]: sns.relplot(x='age' , y='amount', kind="line" , data=df.query('amount <= 350'));
```

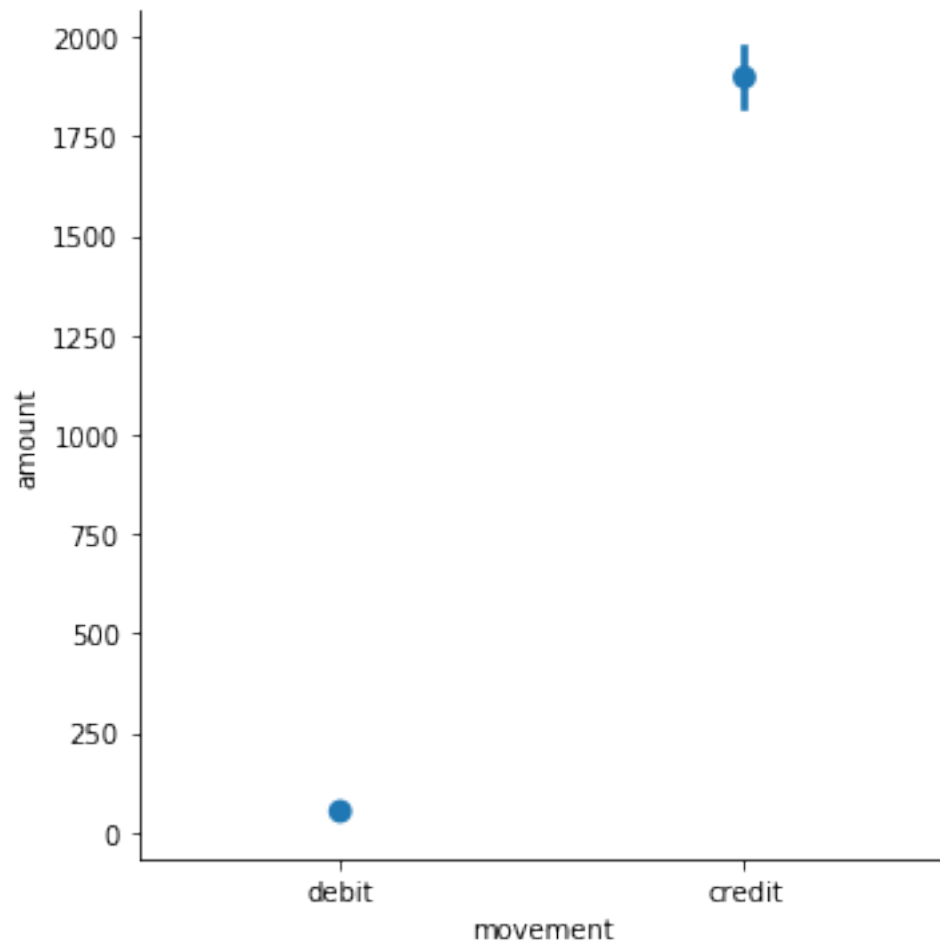


the highest amounts of transactions are made by customers of age between 40 and 45.


8 Amount by transaction type


```
In [52]: sns.catplot(x="movement", y="amount", kind="bar", data=df)
plt.title('average transaction amount per movement')
sns.catplot(x="movement", y="amount", kind="point", data=df, join=False);
```





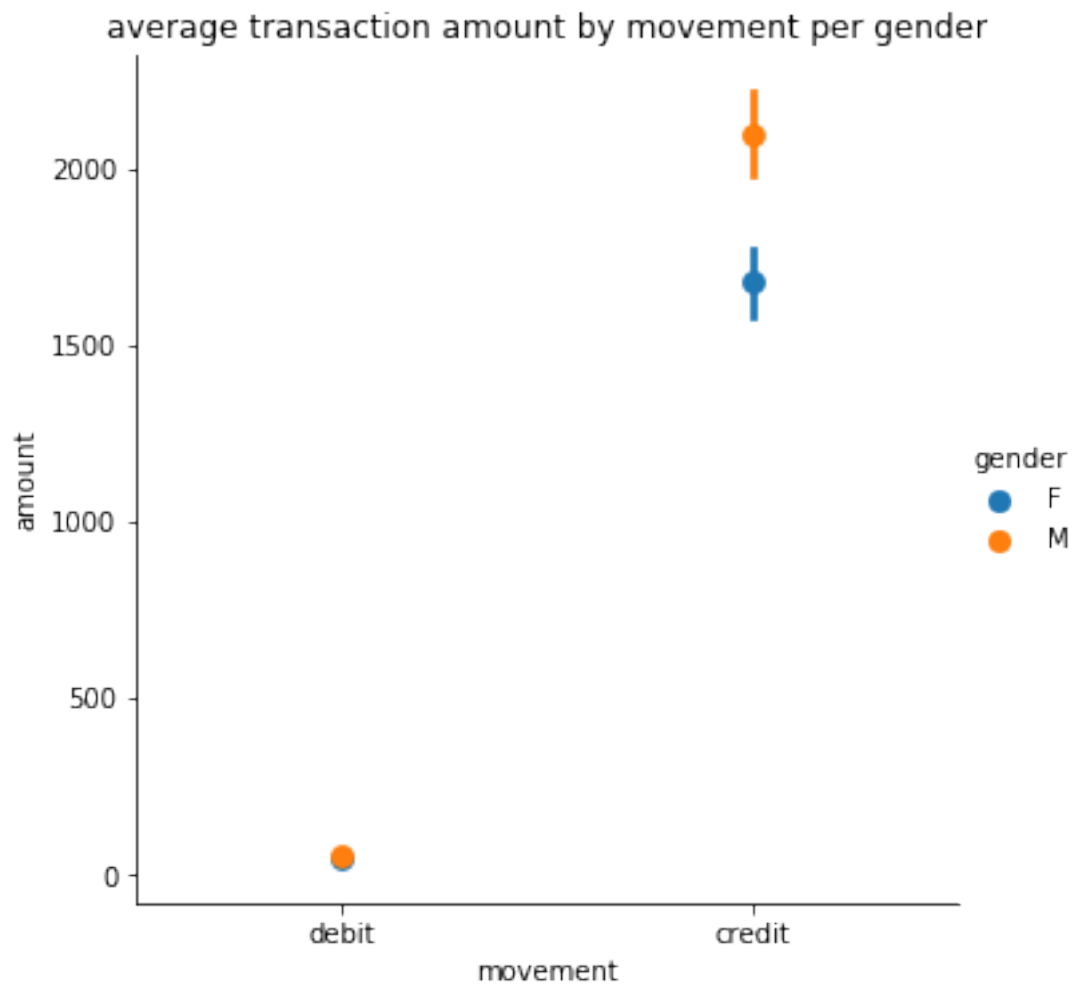
```
In [53]: sns.catplot(x="status", y="amount", kind="bar", data=df.query('amount <= 350'))  
plt.title('average transaction amount per status of transaction')  
sns.catplot(x="status", y="amount", kind="point", data=df.query('amount <= 350'), join
```

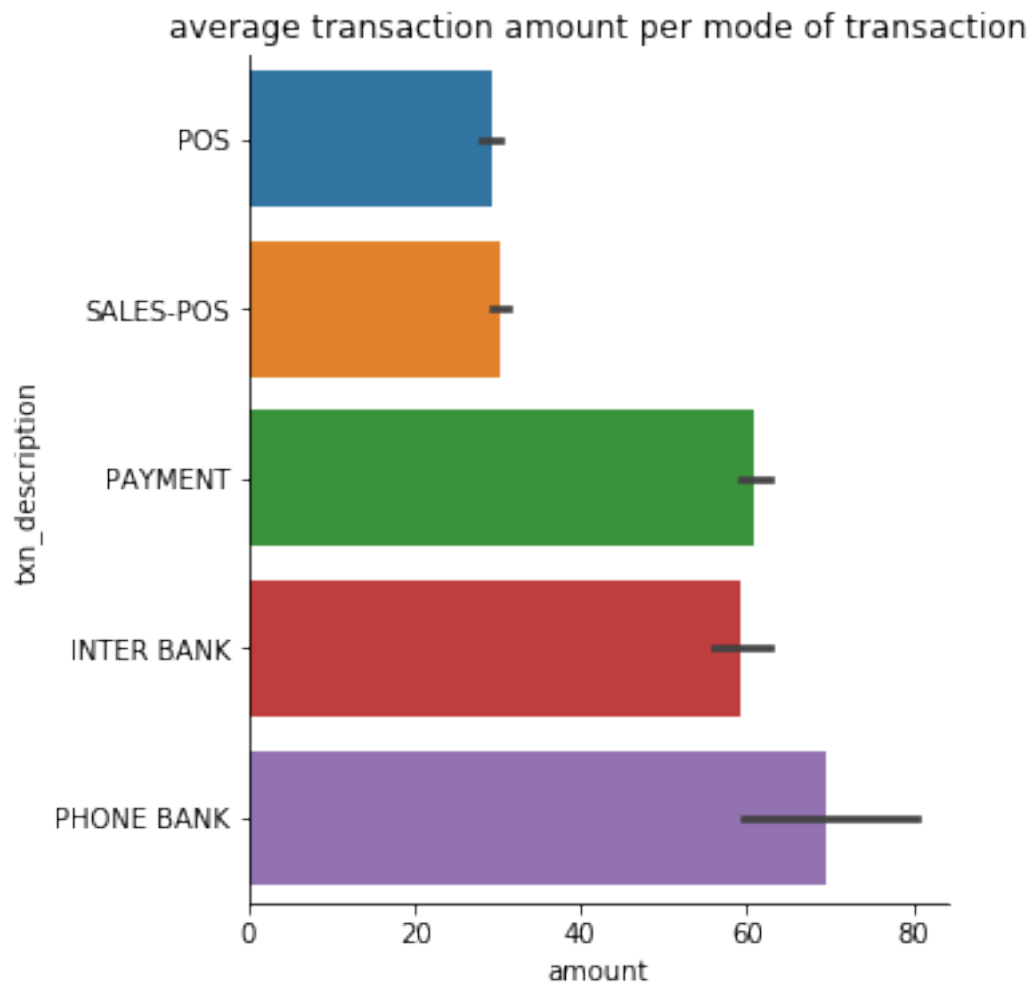


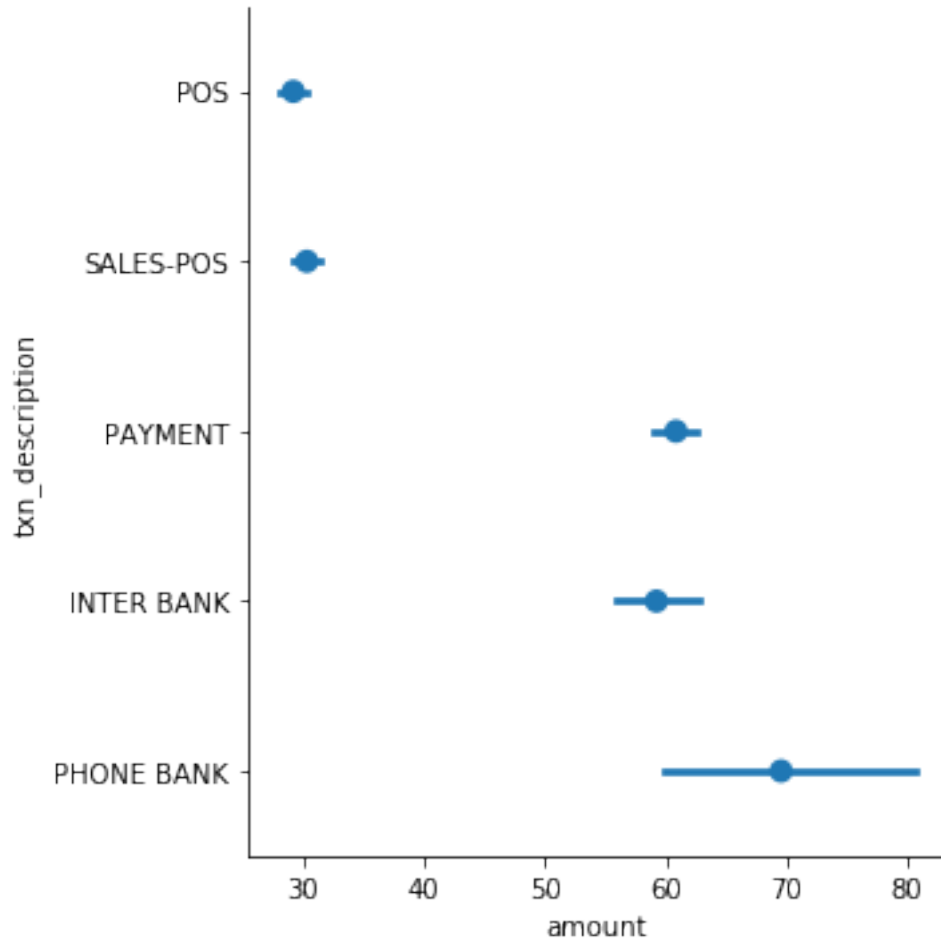
small amount transactions are debit and authorised transactions. high amount transactions are credit and posted transactions.

```
In [54]: sns.catplot(x="movement", y="amount", hue="gender", kind="point", data=df, join=False)
plt.title('average transaction amount by movement per gender');
```




```
In [55]: sns.catplot(y="txn_description", x="amount", kind="bar", data=df.query('amount <= 3500'))  
plt.title('average transaction amount per mode of transaction')  
sns.catplot(y="txn_description", x="amount", kind="point", data=df.query('amount <= 3500'))
```






Phone bank transactions are the ones with the highest amount in transactions. # Amount and Customers


```
In [56]: # customers of the 10 highest transaction amounts
df1 = df.loc[df.amount.sort_values(ascending=False)[:10].index]
sns.catplot(y="customer_id", x="amount", kind="bar", data=df1);
```




```
In [57]: # cutomers with the highest number of transactions
df2 = df.loc[df['customer_id'].isin(df.customer_id.value_counts()[:10].index)]
sns.countplot(y = 'customer_id', data=df2);
```



```
In [58]: # total transactions amounts for the customers with highest number of transactions  
df2.groupby('customer_id')['amount'].sum().plot.bar(alpha = 0.8);
```



```
In [59]: # most spending customers
df.groupby('customer_id')['amount'].sum().sort_values(ascending=False)[:10].plot.bar()
```

```
In [60]: highest_transactions_customers = df1.customer_id.value_counts().index.tolist()
         most_dealing_customers = df2.customer_id.value_counts().index.tolist()
         most_spending_customers = df.groupby('customer_id')['amount'].sum().sort_values(ascending=False).index.tolist()

In [61]: # no customer belongs to top 10 : most spending customers, highest transactions made, and most dealing customers
         set(highest_transactions_customers).intersection(most_dealing_customers, most_spending_customers)

Out[61]: set()

In [62]: # this customer did the 4th highest transaction amount and ranked 9th in the customer_id
         set(highest_transactions_customers).intersection(most_dealing_customers)

Out[62]: {'CUS-51506836'}
```

```
In [63]: # this customer did the highest transaction amount and is the thirteenth most spending customer
         set(highest_transactions_customers).intersection(most_spending_customers)

Out[63]: {'CUS-1816693151'}
```

```
In [64]: # this customer did more transactions than all others, and is the 9th most spending customer
         set(most_dealing_customers).intersection(most_spending_customers)

Out[64]: {'CUS-1196156254'}
```

9 Amount and Balance


```
In [65]: sns.heatmap(df[['amount', 'balance']].corr(), annot = True, fmt = '.3f',  
                  cmap = 'vlag_r', center = 0)  
plt.show()
```



There is no significant relation between the balance and transaction amount.

10 Transaction by Hours

```
In [66]: sns.relplot(x='hour' , y='amount', kind="line" , data=df.query('amount <= 350'));
```



```
In [67]: writer = pd.ExcelWriter('anz_cleaned_data.xlsx',  
                                engine = 'xlsxwriter')
```

```
In [68]: df.to_excel(writer, sheet_name = 'Dataset')
```

```
In [69]: writer.save()  
         writer.close()
```